

In [215]: `%load_ext sql`

The sql extension is already loaded. To reload it, use:
`%reload_ext sql`

In [216]: `%sql mysql://prod:nerd@52.2.153.189/rental_nerd`

Out[216]: u'Connected: prod@rental_nerd'

In [217]: `result = %sql (SELECT \
properties.id as "property_id", \
property_transaction_logs.id as "transaction_log_id", \
properties.*, \
property_transaction_logs.* \
FROM \
properties, \
property_transactions, \
property_transaction_logs \
WHERE \
properties.id = property_transactions.property_id AND \
property_transactions.property_transaction_log_id = property_transaction_logs.id AND \
property_transactions.transaction_type = 'rental')`

`data = result.DataFrame()`

560 rows affected.

In [218]: `result.csv(filename="SQLdump.csv")`

Out[218]: [CSV results \(/files/SQLdump.csv\)](#)

In [219]: `# imports
import pandas as pd
import matplotlib.pyplot as plt
follow the usual sklearn pattern: import, instantiate, fit
from sklearn.linear_model import LinearRegression
import numpy as np

this allows plots to appear directly in the notebook
%matplotlib inline

read data into a DataFrame
data.head()`

Out[219]:

	property_id	transaction_log_id	id	address	neighborhood	bedrooms	bathrooms	sqft	source	origin_url	...	id	price	tra
0	1	1	1	567 Vallejo Street #PH500	San Francisco (North Beach)	3	3	2081	climbsf_renting	http://www.climbsf.com/for-rent/567-vallejo-st...	...	1	12000	op
1	2	2	2	252 Granada Avenue	San Francisco (Ingleside)	2	2	1600	climbsf_renting	http://www.climbsf.com/for-rent/252-granada-ave/	...	2	3950	op
2	3	3	3	460 Valley Street	San Francisco (Noe Valley)	2	2	1446	climbsf_renting	http://www.climbsf.com/for-rent/460-valley-st/	...	3	5400	op
3	4	4	4	333 Fremont Street #705	San Francisco (South Beach)	1	1	0	climbsf_renting	http://www.climbsf.com/for-rent/333-fremont-st...	...	4	3600	op
4	5	5	5	420 Mission Bay Boulevard North #121	San Francisco (Mission Bay)	1	1	980	climbsf_renting	http://www.climbsf.com/for-rent/420-mission-ba...	...	5	3975	op

5 rows × 26 columns

```
In [220]: import datetime

Date_final = [0.1] * len(data)

for x in range(0,len(data)):
    data
    if data["date_closed"][x] is not None :
        # print " row: "+ `x` + ": using date_rented"
        # data.ix['Date_final',x]
        Date_final[x] = data["date_closed"][x]

    elif data["date_listed"][x] is not None :
        # print " row: "+ `x` + ": using date_listed"
        Date_final[x] = data["date_listed"][x]
    else:
        Date_final[x] = data["date_closed"][2]
        print " row: "+ `x` + ": we are screwed"

data['Date'] = pd.to_datetime(Date_final)

data.head()
```

Out[220]:

	property_id	transaction_log_id	id	address	neighborhood	bedrooms	bathrooms	sqft	source	origin_url	...	price	trans
0	1	1	1	567 Vallejo Street #PH500	San Francisco (North Beach)	3	3	2081	climbsf_renting	http://www.climbsf.com/for-rent/567-vallejo-st...	...	12000	open
1	2	2	2	252 Granada Avenue	San Francisco (Ingleside)	2	2	1600	climbsf_renting	http://www.climbsf.com/for-rent/252-granada-ave/	...	3950	open
2	3	3	3	460 Valley Street	San Francisco (Noe Valley)	2	2	1446	climbsf_renting	http://www.climbsf.com/for-rent/460-valley-st/	...	5400	open
3	4	4	4	333 Fremont Street #705	San Francisco (South Beach)	1	1	0	climbsf_renting	http://www.climbsf.com/for-rent/333-fremont-st...	...	3600	open
4	5	5	5	420 Mission Bay Boulevard North #121	San Francisco (Mission Bay)	1	1	980	climbsf_renting	http://www.climbsf.com/for-rent/420-mission-ba...	...	3975	open

5 rows × 27 columns

```
In [221]: # create neighborhoods from lat/long coordinates
import fiona
import shapely as shapely
from shapely.geometry import asShape
```

```
In [227]: shaped_neighborhood = ['None'] * len(data)

with fiona.open('data/Planning_Neighborhoods_4326/planning_hoods_4326.shp') as fiona_collection:
    for hood in fiona_collection:
        print "checking for listings in: " + hood["properties"]["neighborho"]
        # Use Shapely to create the polygon
        shape = asShape( hood['geometry'] )

        for row in range(0,len(data)):
            point = shapely.geometry.Point([data['longitude'][row], data['latitude'][row]]) # longitude, latitude

            if shaped_neighborhood[row] != 'None':
                continue

            if shape.contains(point):
                #print `row` + ": Found " + data.address[row] + " in hood " + hood["properties"]["nbrhood"]
                shaped_neighborhood[row] = hood["properties"]["neighborho"]

data['shaped_neighborhood'] = shaped_neighborhood
data.head()
```

checking for listings in: Seacliff
 checking for listings in: Haight Ashbury
 checking for listings in: Outer Mission
 checking for listings in: Russian Hill
 checking for listings in: Noe Valley
 checking for listings in: Inner Sunset
 checking for listings in: Downtown/Civic Center
 checking for listings in: Diamond Heights
 checking for listings in: Treasure Island/YBI
 checking for listings in: Lakeshore
 checking for listings in: Outer Richmond
 checking for listings in: Crocker Amazon
 checking for listings in: Excelsior
 checking for listings in: Parkside
 checking for listings in: Financial District
 checking for listings in: Ocean View
 checking for listings in: Mission
 checking for listings in: West of Twin Peaks
 checking for listings in: Inner Richmond
 checking for listings in: Marina
 checking for listings in: Bayview
 checking for listings in: Visitacion Valley
 checking for listings in: Pacific Heights
 checking for listings in: Presidio
 checking for listings in: Nob Hill
 checking for listings in: Outer Sunset
 checking for listings in: Western Addition
 checking for listings in: Golden Gate Park
 checking for listings in: Presidio Heights
 checking for listings in: South of Market
 checking for listings in: Glen Park
 checking for listings in: Potrero Hill
 checking for listings in: Castro/Upper Market
 checking for listings in: Twin Peaks
 checking for listings in: Bernal Heights
 checking for listings in: Chinatown
 checking for listings in: North Beach

Out[227]:

	property_id	transaction_log_id	id	address	neighborhood	bedrooms	bathrooms	sqft	source	origin_url	...	transaction_s
0	1	1	1	567 Vallejo Street #PH500	San Francisco (North Beach)	3	3	2081	climbsf_renting	http://www.climbsf.com/for-rent/567-vallejo-st...	...	open
1	2	2	2	252 Granada Avenue	San Francisco (Ingleside)	2	2	1600	climbsf_renting	http://www.climbsf.com/for-rent/252-granada-ave/	...	open
2	3	3	3	460 Valley Street	San Francisco (Noe Valley)	2	2	1446	climbsf_renting	http://www.climbsf.com/for-rent/460-valley-st/	...	open
3	4	4	4	333 Fremont Street #705	San Francisco (South Beach)	1	1	0	climbsf_renting	http://www.climbsf.com/for-rent/333-fremont-st...	...	open
4	5	5	5	420 Mission Bay Boulevard North #121	San Francisco (Mission Bay)	1	1	980	climbsf_renting	http://www.climbsf.com/for-rent/420-mission-ba...	...	open

5 rows × 28 columns

```
In [228]: # filter out any outliers, defined as rent >$10k or >2,500 sq ft, or not in SF

print "Entries before filter: " + `len(data)`
data = data[(data.shaped_neighborhood != 'None') & (data.sqft <= 2500) & (data.price <= 8000) & (data.price != 0) & (data.
bedrooms <= 4) & (data.bathrooms <= 3) & (data.sqft != 0)]

# filter out listings over one month old

print "Entries after filter: " + `len(data)`

Entries before filter: 560
Entries after filter: 304
```

```
In [229]: # create year dummy variables (because date isn't very intuitive variable)
data["Year"] = pd.DatetimeIndex(data["Date"]).to_period('Y')

# create dummy variables using get_dummies, then exclude the first dummy column
year_dummies = pd.get_dummies(data.Year, prefix='Year').iloc[:, :-1]

# print out baseline neighborhood
base_area = pd.get_dummies(data.shaped_neighborhood, prefix='neighborhood').iloc[:, 0:1].columns[0]
print('Base neighborhood: %s' % base_area)

# create dummy variables using get_dummies, then exclude the first dummy column
area_dummies = pd.get_dummies(data.shaped_neighborhood, prefix='neighborhood').iloc[:, 1:]

# concatenate the dummy variable columns onto the original DataFrame (axis=0 means rows, axis=1 means columns)
data = pd.concat([data, area_dummies, year_dummies], axis=1)

data.head()
```

Base neighborhood: neighborhood_Bayview

```
Out[229]:
```

	property_id	transaction_log_id	id	address	neighborhood	bedrooms	bathrooms	sqft	source	origin_url	...	neighborho Hill
1	2	2	2	252 Granada Avenue	San Francisco (Ingleside)	2	2	1600	climbsf_renting	http://www.climbsf.com/for-rent/252-granada-ave/	...	0
2	3	3	3	460 Valley Street	San Francisco (Noe Valley)	2	2	1446	climbsf_renting	http://www.climbsf.com/for-rent/460-valley-st/	...	0
4	5	5	5	420 Mission Bay Boulevard North #121	San Francisco (Mission Bay)	1	1	980	climbsf_renting	http://www.climbsf.com/for-rent/420-mission-ba...	...	0
7	8	8	8	1160 Mission Street #1112	San Francisco (SOMA)	1	1	664	climbsf_renting	http://www.climbsf.com/for-rent/1160-mission-s...	...	0
11	12	12	12	655 26th Avenue	San Francisco (Central Richmond)	2	1	1300	climbsf_renting	http://www.climbsf.com/for-rent/655-26th-ave/	...	0

5 rows x 62 columns

```

In [230]: # FACTORING BY YEAR AND NEIGHBORHOOD
# Thesis: Neighborhoods influence valuations as a multiplier, rather than a constant.
# a square foot in SOMA is worth more than a square foot in Portrero by X%
# New model will look like this:
# Price = B_1 x (SOMA Coeff * Year Coeff * Sqft) + intercept
# $3,900 = B_1 x (1.20% * 1.15% * 2,023 sqft) + intercept
# where B_1 represents the price per square foot in base year and base neighborhood
# I will ignore intercepts for now FIXME
# calculate the coefficients for the following matrix and save them for later regressions
#           SOMA    Mission    Portrero    Intercept
# Price/SQFT    $1.23    $0.59    $0.88    $_.__

# create Price per square foot

price_per_foot = data.price / data.sqft
price_per_foot.name = 'price_per_foot'
data = pd.concat([data, price_per_foot], axis=1)

data.head()

```



Out[230]:

	property_id	transaction_log_id	id	address	neighborhood	bedrooms	bathrooms	sqft	source	origin_url	...	neighborho of Market
1	2	2	2	252 Granada Avenue	San Francisco (Ingleside)	2	2	1600	climbsf_renting	http://www.climbsf.com/for- rent/252-granada-ave/	...	0
2	3	3	3	460 Valley Street	San Francisco (Noe Valley)	2	2	1446	climbsf_renting	http://www.climbsf.com/for- rent/460-valley-st/	...	0
4	5	5	5	420 Mission Bay Boulevard North #121	San Francisco (Mission Bay)	1	1	980	climbsf_renting	http://www.climbsf.com/for- rent/420-mission-ba...	...	1
7	8	8	8	1160 Mission Street #1112	San Francisco (SOMA)	1	1	664	climbsf_renting	http://www.climbsf.com/for- rent/1160-mission-s...	...	1
11	12	12	12	655 26th Avenue	San Francisco (Central Richmond)	2	1	1300	climbsf_renting	http://www.climbsf.com/for- rent/655-26th-ave/	...	0

5 rows × 63 columns

```
In [231]: # visualize the relationship between the features and the response using scatterplots
data.plot(kind='scatter', x='elevation', y='price')
data.plot(kind='scatter', x='elevation', y='price_per_foot')
```

```
Out[231]: <matplotlib.axes._subplots.AxesSubplot at 0x110e87890>
```



```
In [232]: class ListTable(list):
    """ Overridden list class which takes a 2-dimensional list of
    the form [[1,2,3],[4,5,6]], and renders an HTML Table in
    IPython Notebook. """

    def _repr_html_(self):
        html = ["<table>"]
        for row in self:
            html.append("<tr>")

            for col in row:
                html.append("<td>{0}</td>".format(col))

            html.append("</tr>")
        html.append("</table>")
        return ''.join(html)
```

```
In [233]: feature_cols = area_dummies.columns

X = data[feature_cols]
y = data.price_per_foot

# instantiate, fit
lm = LinearRegression()
lm.fit(X, y)

# print coefficients
# The mean square error
print("Residual sum of squares: %.2f"
      % np.mean((lm.predict(X) - y) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % lm.score(X, y))

# print raw results
print("Base area is %s: $%.2f" % (base_area, lm.intercept_))

zip(feature_cols, lm.coef_)

table = ListTable()

dtype = [('Neighborhood', 'S100'), ('$ per square', float)]

# round to pennies
round_coef = map(round, lm.coef_, [2]*len(lm.coef_))
x = np.array(zip(feature_cols, round_coef), dtype=dtype)
x.T
x = np.sort(x, axis=0, order='$ per square')

table.append(['Neighborhood', '$ per square (+/-)'])
for i in x:
    table.append(i)

table
```


Residual sum of squares: 1.04
 Variance score: 0.29
 Base area is neighborhood_Bayview: \$3.15

Out[233]:

Neighborhood	\$ per square (+/-)
neighborhood_Visitacion Valley	-0.58
neighborhood_Glen Park	-0.42
neighborhood_Bernal Heights	-0.36
neighborhood_Diamond Heights	-0.34
neighborhood_Lakeshore	-0.28
neighborhood_Excelsior	-0.21
neighborhood_Ocean View	0.06
neighborhood_Outer Richmond	0.13
neighborhood_Haight Ashbury	0.14
neighborhood_Parkside	0.37
neighborhood_West of Twin Peaks	0.38
neighborhood_Outer Sunset	0.4
neighborhood_Potrero Hill	1.02
neighborhood_Noë Valley	1.07
neighborhood_Marina	1.12
neighborhood_South of Market	1.13
neighborhood_Inner Richmond	1.24
neighborhood_Downtown/Civic Center	1.26
neighborhood_North Beach	1.27
neighborhood_Nob Hill	1.32
neighborhood_Pacific Heights	1.52
neighborhood_Mission	1.62
neighborhood_Financial District	1.65
neighborhood_Castro/Upper Market	1.68
neighborhood_Inner Sunset	1.94
neighborhood_Russian Hill	2.48
neighborhood_Western Addition	2.53
neighborhood_Chinatown	5.19

```

In [234]: full_price = [lm.intercept_] * len(lm.coef_)
full_price += lm.coef_

area_price_per_foot = dict(zip(feature_cols,full_price))
area_price_per_foot[base_area] = lm.intercept_

dtype = [('Neighborhood', 'S100'), ('$ per sqft', float)]

# round to pennies
round_coef = map(round,full_price,[2]*len(full_price))
x = np.array(zip(feature_cols, full_price),dtype=dtype)
x.T
x = np.sort(x,axis=0,order='$ per sqft')

table = ListTable()

table.append(['Neighborhood','$ per sqft'])
for i in x:
    table.append(i)

table

```

Out[234]:

Neighborhood	\$ per sqft
neighborhood_Visitacion Valley	2.56314257913
neighborhood_Glen Park	2.72727272727
neighborhood_Bernal Heights	2.78200061463
neighborhood_Diamond Heights	2.8085106383
neighborhood_Lakeshore	2.86666666667
neighborhood_Excelsior	2.93253968254
neighborhood_Ocean View	3.21056547619
neighborhood_Outer Richmond	3.27613474488
neighborhood_Haight Ashbury	3.28587304169
neighborhood_Parkside	3.51137487636
neighborhood_West of Twin Peaks	3.53093930792
neighborhood_Outer Sunset	3.54200988468
neighborhood_Potrero Hill	4.16518162307
neighborhood_Noel Valley	4.2183338668
neighborhood_Marina	4.27103404056
neighborhood_South of Market	4.27811516862
neighborhood_Inner Richmond	4.38912151969
neighborhood_Downtown/Civic Center	4.40392841058
neighborhood_North Beach	4.42100898552
neighborhood_Nob Hill	4.46503589805
neighborhood_Pacific Heights	4.66270219935
neighborhood_Mission	4.76198374942
neighborhood_Financial District	4.79539277753
neighborhood_Castro/Upper Market	4.8298805149
neighborhood_Inner Sunset	5.08241758242
neighborhood_Russian Hill	5.62111377504
neighborhood_Western Addition	5.67123941357
neighborhood_Chinatown	8.33333333333

```

In [235]: # calculate the multipliers for each neighborhood relative to base area
# SOMA_mult = SOMA_per_foot / Base_per_foot

area_mults = [lm.intercept_] * len(lm.coef_)
area_mults = full_price / area_mults - [1]*len(lm.coef_)

dtype = [('Neighborhood', 'S100'), ('Multiplier', float)]

# round to pennies
round_coef = map(round,area_mults,[2]*len(area_mults))
x = np.array(zip(feature_cols, area_mults),dtype=dtype)
x.T
x = np.sort(x,axis=0,order='Multiplier')

table = ListTable()

table.append(['Neighborhood','Multiplier'])
table.append([base_area,0])
for i in x:
    table.append(i)

table

```

Out[235]:

Neighborhood	Multiplier
neighborhood_Bayview	0
neighborhood_Visitacion Valley	-0.185286214946
neighborhood_Glen Park	-0.133116236059
neighborhood_Bernal Heights	-0.115720573165
neighborhood_Diamond Heights	-0.107294166495
neighborhood_Lakeshore	-0.08880884368
neighborhood_Excelsior	-0.0678706194346
neighborhood_Ocean View	0.0205019309387
neighborhood_Outer Richmond	0.0413436069017
neighborhood_Haight Ashbury	0.04443899641
neighborhood_Parkside	0.11611641879
neighborhood_West of Twin Peaks	0.122335117748
neighborhood_Outer Sunset	0.125853982273
neighborhood_Potrero Hill	0.323933718397
neighborhood_Noel Valley	0.340828551337
neighborhood_Marina	0.357579690501
neighborhood_South of Market	0.359830479313
neighborhood_Inner Richmond	0.395114667242
neighborhood_Downtown/Civic Center	0.399821146787
neighborhood_North Beach	0.405250333588
neighborhood_Nob Hill	0.419244612659
neighborhood_Pacific Heights	0.482074305325
neighborhood_Mission	0.513631678725
neighborhood_Financial District	0.524250984873
neighborhood_Castro/Upper Market	0.535213166719
neighborhood_Inner Sunset	0.61548393738
neighborhood_Russian Hill	0.786712497844
neighborhood_Western Addition	0.802645301983
neighborhood_Chinatown	1.64881150093

In [236]: *# calculate the adjusted Sqft (Sqft * Area_mult) for the dataset and add it as a new column to data*

for each property, multiplier is sum of array [area_dummies] x [area_mults]

```
t = data[area_dummies.columns] * area_mults
t = t.T.sum()
```

```
t.name = 'area_multiplier'
t = t + 1
data = pd.concat([data, t], axis=1)
```

```
adj_sqft = data.sqft * t
adj_sqft.name = 'area_adj_sqft'
data = pd.concat([data, adj_sqft], axis=1)
```

```
data.head()
```

Out[236]:

	property_id	transaction_log_id	id	address	neighborhood	bedrooms	bathrooms	sqft	source	origin_url	...	neighborho of Twin Pea
1	2	2	2	252 Granada Avenue	San Francisco (Ingleside)	2	2	1600	climbsf_renting	http://www.climbsf.com/for- rent/252-granada-ave/	...	0
2	3	3	3	460 Valley Street	San Francisco (Noe Valley)	2	2	1446	climbsf_renting	http://www.climbsf.com/for- rent/460-valley-st/	...	0
4	5	5	5	420 Mission Bay Boulevard North #121	San Francisco (Mission Bay)	1	1	980	climbsf_renting	http://www.climbsf.com/for- rent/420-mission-ba...	...	0
7	8	8	8	1160 Mission Street #1112	San Francisco (SOMA)	1	1	664	climbsf_renting	http://www.climbsf.com/for- rent/1160-mission-s...	...	0
11	12	12	12	655 26th Avenue	San Francisco (Central Richmond)	2	1	1300	climbsf_renting	http://www.climbsf.com/for- rent/655-26th-ave/	...	0

5 rows × 65 columns

```

In [237]: # run the regression based on area_adj_sqft rather than sqft

# create X and y
feature_cols = [data.area_adj_sqft.name]

X = data[feature_cols]
y = data.price

# instantiate, fit
lm = LinearRegression()
lm.fit(X, y)

# print coefficients
print("Intercept: %.2f" % lm.intercept_)

# The mean square error
print("Residual sum of squares: %.2f"
      % np.mean((lm.predict(X) - y) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % lm.score(X, y))
zip(feature_cols, lm.coef_)

# calculate predictions for the data set and plot errors
predictions = lm.predict(X)
errors = predictions-y
errors.name = 'Error'

# visualize the relationship between the features and the response using scatterplots
errors.sort()
errors.plot(kind='bar').get_xaxis().set_ticks([])


```

```

Intercept: 1616.99
Residual sum of squares: 611848.95
Variance score: 0.66

```

Out[237]: []



In [238]: feature_cols = year_dummies.columns

```
X = data[feature_cols]
y = data.price_per_foot

# instantiate, fit
lm = LinearRegression()
lm.fit(X, y)

# print coefficients
# The mean square error
print("Residual sum of squares: %.2f"
      % np.mean((lm.predict(X) - y) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % lm.score(X, y))

# print raw results
print lm.intercept_

zip(feature_cols, lm.coef_)
```

Residual sum of squares: 1.32
Variance score: 0.11
4.79633313103

Out[238]: [(u'Year_1969', -1.9410031817959141),
(u'Year_2011', -0.62577229712902971),
(u'Year_2012', -1.3700383103173595),
(u'Year_2013', -0.98855679368271021),
(u'Year_2014', -0.60179239513716498)]

In [239]: full_price = [lm.intercept_] * len(lm.coef_)
full_price += lm.coef_

year_price_per_foot = dict(zip(feature_cols, full_price))
year_price_per_foot[base_area] = lm.intercept_

print year_price_per_foot

{u'Year_1969': 2.8553299492385782, u'Year_2012': 3.4262948207171329, u'Year_2013': 3.8077763373517821, u'neighborhood_Bayview': 4.7963331310344923, u'Year_2011': 4.1705608339054629, u'Year_2014': 4.194540735897327}

In [240]: # calculate the multipliers for each year relative to base year
2014_mult = 2014_per_foot / 2015_per_foot

year_mults = [lm.intercept_] * len(lm.coef_)
year_mults = full_price / year_mults - [1]*len(lm.coef_)

zip(feature_cols, year_mults)

Out[240]: [(u'Year_1969', -0.40468481416287083),
(u'Year_2011', -0.13046889780861826),
(u'Year_2012', -0.28564285942787804),
(u'Year_2013', -0.2061067833854765),
(u'Year_2014', -0.12546926551937987)]

In [241]: *# calculate the adjusted Sqft (Sqft * Year_mult) for the dataset and add it as a new column to data*

for each property, multiplier is sum of array [year_dummies] x [year_mults]

```
t = data[year_dummies.columns] * year_mults
t = t.T.sum()
```

```
t.name = 'year_multiplier'
```

```
t = t + 1
```

```
data = pd.concat([data, t], axis=1)
```

```
year_adj_sqft = data.area_adj_sqft * t
```

```
year_adj_sqft.name = 'adj_sqft'
```

```
data = pd.concat([data, year_adj_sqft], axis=1)
```

```
data.head()
```

Out[241]:

	property_id	transaction_log_id	id	address	neighborhood	bedrooms	bathrooms	sqft	source	origin_url	...	Year_1969
1	2	2	2	252 Granada Avenue	San Francisco (Ingleside)	2	2	1600	climbsf_renting	http://www.climbsf.com/for-rent/252-granada-ave/	...	0
2	3	3	3	460 Valley Street	San Francisco (Noe Valley)	2	2	1446	climbsf_renting	http://www.climbsf.com/for-rent/460-valley-st/	...	0
4	5	5	5	420 Mission Bay Boulevard North #121	San Francisco (Mission Bay)	1	1	980	climbsf_renting	http://www.climbsf.com/for-rent/420-mission-ba...	...	0
7	8	8	8	1160 Mission Street #1112	San Francisco (SOMA)	1	1	664	climbsf_renting	http://www.climbsf.com/for-rent/1160-mission-s...	...	0
11	12	12	12	655 26th Avenue	San Francisco (Central Richmond)	2	1	1300	climbsf_renting	http://www.climbsf.com/for-rent/655-26th-ave/	...	0

5 rows × 67 columns

```

In [242]: # run the regression based on year_and_area_adj_sqft rather than area_adj_sqft

# create X and y
feature_cols = ['adj_sqft']

X = data[feature_cols]
y = data.price

# instantiate, fit
lm = LinearRegression()
lm.fit(X, y)

# print coefficients
print lm.intercept_
# The mean square error
print("Residual sum of squares: %.2f"
      % np.mean((lm.predict(X) - y) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % lm.score(X, y))
print zip(feature_cols, lm.coef_)

# calculate predictions for the data set and plot errors
predictions = lm.predict(X)
errors = predictions-y
errors.name = 'Error'

# visualize the relationship between the features and the response using scatterplots
errors.sort(inplace=True)
errors.plot(kind='bar').get_xaxis().set_ticks([])

errors.tail(10)

```

```


1507.8595447
Residual sum of squares: 552478.52
Variance score: 0.69
[('adj_sqft', 2.1307024315775638)]

```

```

Out[242]: 202    1384.734919
194    1417.817265
318    1441.690199
313    1457.964199
108    1463.663204
294    1476.202741
236    1481.893854
455    1827.549514
60     2224.440816
523    2608.505581
Name: Error, dtype: float64

```




```

In [243]: # create X and y
feature_cols = ['adj_sqft', 'bedrooms', 'bathrooms']

X = data[feature_cols]
y = data.price

# instantiate, fit
lm = LinearRegression()
lm.fit(X, y)

# print coefficients
print("Intercept: %.2f" % lm.intercept_)
# The mean square error
print("Residual sum of squares: %.2f"
      % np.mean((lm.predict(X) - y) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % lm.score(X, y))
print(zip(feature_cols, lm.coef_))


# calculate predictions for the data set and plot errors
predictions = lm.predict(X)
errors = predictions - y
errors.name = 'Error'

# visualize the relationship between the features and the response using scatterplots
errors.sort()
errors.plot(kind='bar').get_xaxis().set_ticks([])

Intercept: 1269.34
Residual sum of squares: 509648.11
Variance score: 0.71
[('adj_sqft', 1.7729741412548015), ('bedrooms', 155.64516216444514), ('bathrooms', 323.83523244860413)]

```

Out[243]: []



In [244]: *# show errors by neighborhood to see if there are any neighborhoods with funky differences*

```
hooderrors = data[['neighborhood']]

errors = predictions-y
errors.name = 'Error'

hooderrors = pd.concat([hooderrors,errors.abs()],axis=1)

hood_group = hooderrors.groupby('neighborhood')

import numpy
def median(lst):
    return numpy.median(numpy.array(lst))

error_avg = hood_group.median()
error_avg.sort(columns='Error',ascending=False).plot(kind='bar')

# show errors by year to see if there are any years with funky differences

yearerrors = data[['Year']]

yearerrors = pd.concat([yearerrors,errors.abs()],axis=1)

year_group = yearerrors.groupby('Year')
error_avg = year_group.mean()
error_avg.sort(columns='Error',ascending=False).plot(kind='bar')




# show errors by source to see if there are any sources have noisy data

srcerrors = data[['source']]

srcerrors = pd.concat([srcerrors,errors.abs()],axis=1)

src_group = srcerrors.groupby('source')
error_avg = src_group.mean()
error_avg.sort(columns='Error',ascending=False).plot(kind='bar')
```

Out[244]: <matplotlib.axes._subplots.AxesSubplot at 0x111d8ac50>



```
In [245]: import csv

table = ListTable()

dtype = [('Effect', 'S100'), ('Coefficient', float)]

# round to pennies
round_coef = map(round, lm.coef_, [6]*len(lm.coef_))
x = np.array(zip(feature_cols, round_coef), dtype=dtype)
x.T
print zip(feature_cols, lm.coef_)
#x = np.sort(x, axis=0, order='Coefficient')

with open('model_features_v1.csv', 'wb') as csvfile:
    modelwriter = csv.writer(csvfile, delimiter=',', quotechar='|', quoting=csv.QUOTE_MINIMAL)

    header = ['Effect', 'Coefficient']
    table.append(header)
    modelwriter.writerow(header)
    for i in x:
        table.append(i)
        modelwriter.writerow(i)

    table.append(['base_rent', lm.intercept_])

    modelwriter.writerow(['base_rent', lm.intercept_])

table

[('adj_sqft', 1.7729741412548015), ('bedrooms', 155.64516216444514), ('bathrooms', 323.83523244860413)]
```

```
Out[245]:
```

Effect	Coefficient
adj_sqft	1.772974
bedrooms	155.645162
bathrooms	323.835232
base_rent	1269.34366491

```

In [246]: table = ListTable()

dtype = [('Effect', 'S100'), ('Coefficient', float)]

# round to pennies
round_coef = map(round, (area_mults + [1]*len(area_mults)), [6]*len(area_mults))
x = np.array(zip(area_dummies.columns, round_coef), dtype=dtype)
x.T
x = np.sort(x, axis=0, order='Coefficient')

with open('model_hoods_v1.csv', 'wb') as csvfile:
    hoodwriter = csv.writer(csvfile, delimiter=',', quotechar='|', quoting=csv.QUOTE_MINIMAL)

    header = ['Neighborhood', 'Multiplier']
    table.append(header)
    hoodwriter.writerow(header)

    for i in x:
        table.append(i)
        hoodwriter.writerow(i)

    lastrow = [base_area, 1]
    table.append(lastrow)
    hoodwriter.writerow(lastrow)

table

```

Out[246]:

Neighborhood	Multiplier
neighborhood_Visitacion Valley	0.814714
neighborhood_Glen Park	0.866884
neighborhood_Bernal Heights	0.884279
neighborhood_Diamond Heights	0.892706
neighborhood_Lakeshore	0.911191
neighborhood_Excelsior	0.932129
neighborhood_Ocean View	1.020502
neighborhood_Outer Richmond	1.041344
neighborhood_Haight Ashbury	1.044439
neighborhood_Parkside	1.116116
neighborhood_West of Twin Peaks	1.122335
neighborhood_Outer Sunset	1.125854
neighborhood_Potrero Hill	1.323934
neighborhood_Noë Valley	1.340829
neighborhood_Marina	1.35758
neighborhood_South of Market	1.35983
neighborhood_Inner Richmond	1.395115
neighborhood_Downtown/Civic Center	1.399821
neighborhood_North Beach	1.40525
neighborhood_Nob Hill	1.419245
neighborhood_Pacific Heights	1.482074
neighborhood_Mission	1.513632
neighborhood_Financial District	1.524251
neighborhood_Castro/Upper Market	1.535213
neighborhood_Inner Sunset	1.615484
neighborhood_Russian Hill	1.786712
neighborhood_Western Addition	1.802645
neighborhood_Chinatown	2.648812
neighborhood_Bayview	1

In [247]: *# show negative errors meaning we expected rents to be higher*

```
error = predictions-y
error.name = 'error'

data = pd.concat([data,error,pd.DataFrame(predictions,columns=['predicted_price']),axis=1)

data.head()
```

Out[247]:

	property_id	transaction_log_id	id	address	neighborhood	bedrooms	bathrooms	sqft	source	origin_url	...	Year_2012
0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN
1	2	2	2	252 Granada Avenue	San Francisco (Ingleside)	2	2	1600	climbsf_renting	http://www.climbsf.com/for-rent/252-granada-ave/	...	0
2	3	3	3	460 Valley Street	San Francisco (Noe Valley)	2	2	1446	climbsf_renting	http://www.climbsf.com/for-rent/460-valley-st/	...	0
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN
4	5	5	5	420 Mission Bay Boulevard North #121	San Francisco (Mission Bay)	1	1	980	climbsf_renting	http://www.climbsf.com/for-rent/420-mission-ba...	...	0

5 rows × 69 columns

```
# filter out overshoot error
overshoot = data[(data.error <= -500)]
columns = data.columns - ['error','latitude', 'longitude', 'address', 'origin_url','price','neighborhood']
overshoot = data.drop(columns,1)
overshoot.sort('error',ascending=True,inplace=True)
overshoot.head(30)
```

Out[248]:

	address	neighborhood	origin_url	latitude	longitude	price	error
546	301 Main St UNIT 35A, San Francisco, CA 94105	South Beach	http://www.zillow.com/homedetails/301-Main-St-...	37.7894	-122.391	7950	-2664.618204
233	338 Spear Street #39E	San Francisco (South Beach)	http://www.climbsf.com/for-rent/338-spear-st-39e/	37.7894	-122.391	7975	-2255.490601
517	20th St San Francisco, CA 94110	None	http://www.zillow.com/homedetails/20th-St-San-...	37.7588	-122.416	6200	-2035.909097
158	338 Spear Street #39A	San Francisco (South Beach)	http://www.climbsf.com/for-rent/338-spear-st-39a/	37.7894	-122.391	6700	-1983.453547
273	301 Mission Street #29F	San Francisco (South Beach)	http://www.climbsf.com/for-rent/301-mission-st...	37.7905	-122.396	7975	-1962.920623
89	88 King Street #904	San Francisco (South Beach)	http://www.climbsf.com/for-rent/88-king-st-904/	37.7807	-122.389	6250	-1895.205653
434	748 Bay St, San Francisco, CA 94109	Russian Hill	http://www.zillow.com/homedetails/748-Bay-St-S...	37.8049	-122.419	7500	-1771.282008
299	401 Harrison Street #3803	San Francisco (Rincon Hill)	http://www.climbsf.com/for-rent/401-harrison-s...	37.7864	-122.392	7225	-1768.441160
232	2560 Vallejo Street	San Francisco (Pacific Heights)	http://www.climbsf.com/for-rent/2560-vallejo-st/	37.7950	-122.439	7050	-1678.668048
381	301 Main Street #35F	San Francisco (South Beach)	http://www.climbsf.com/for-rent/301-main-st-35f/	37.7894	-122.391	7000	-1582.016268
525	20th St San Francisco, CA 94114	None	http://www.zillow.com/homedetails/20th-St-San-...	37.7578	-122.432	5700	-1481.921519
382	480 Mission Bay Boulevard North #PH1606	San Francisco (Mission Bay)	http://www.climbsf.com/for-rent/480-mission-ba...	37.7731	-122.393	7500	-1472.047366
283	234 Grand View Avenue	San Francisco (Noe Valley)	http://www.climbsf.com/for-rent/234-grand-view...	37.7545	-122.441	7300	-1419.202673
411	Vallejo St San Francisco, CA 94133	None	http://www.zillow.com/homedetails/Vallejo-	37.7985	-122.410	4000	-1229.429094

			St-S...				
293	425 1st Street #3402	San Francisco (Rincon Hill)	http://www.climbsf.com/for-rent/425-1st-st-3402/	37.7858	-122.392	6600	-1215.769488
203	461 2nd St. #557T	San Francisco (South Beach)	http://www.climbsf.com/for-rent/461-2nd-st-557t/	37.7838	-122.394	6750	-1180.673838
459	Tehama St San Francisco, CA 94103	None	http://www.zillow.com/homedetails/Tehama-St-Sa...	37.7793	-122.407	6000	-1167.299698
457	Lombard St San Francisco, CA 94133	None	http://www.zillow.com/homedetails/Lombard-St-S...	37.8021	-122.419	6700	-1152.566463
119	1837 Jefferson Street	San Francisco (Marina)	http://www.climbsf.com/for-rent/1837-jefferson...	37.8045	-122.443	6200	-1140.531105
357	296 Francisco Street	San Francisco (Telegraph Hill)	http://www.climbsf.com/for-rent/296-francisco-st/	37.8053	-122.410	5475	-1111.447417
316	35 Dolores Street #410	San Francisco (Mission Dolores)	http://www.climbsf.com/for-rent/35-dolores-st-...	37.7686	-122.427	6050	-1099.802300
204	1839 Jefferson Street	San Francisco (Marina)	http://www.climbsf.com/for-rent/1839-jefferson...	37.8048	-122.443	6400	-1098.461283
134	301 Main Street #25E	San Francisco (South Beach)	http://www.climbsf.com/for-rent/301-main-st-25e/	37.7894	-122.391	5800	-1083.730601
282	301 Main Street #5C	San Francisco (South Beach)	http://www.climbsf.com/for-rent/301-main-st-5c/	37.7894	-122.391	7000	-1018.663679
405	Vallejo St San Francisco, CA 94123	None	http://www.zillow.com/homedetails/Vallejo-St-S...	37.7952	-122.435	4200	-998.230614
109	229 Brannan Street #12J	San Francisco (South Beach)	http://www.climbsf.com/for-rent/229-brannan-st...	37.7826	-122.390	5950	-994.199509
113	301 Mission Street #701	San Francisco (SOMA)	http://www.climbsf.com/for-rent/301-mission-st...	37.7905	-122.396	7400	-977.315886
430	501 Beale St, San Francisco, CA 94105	South Beach	http://www.zillow.com/homedetails/501-Beale-St...	37.7863	-122.389	6000	-975.000185
174	301 Main Street #14F	San Francisco (South Beach)	http://www.climbsf.com/for-rent/301-main-st-14f/	37.7894	-122.391	5950	-965.958102
72	235 Berry Street #107	San Francisco (Mission Bay)	http://www.climbsf.com/for-rent/235-berry-st-1...	37.7749	-122.394	7500	-964.150538

```

In [249]: data = data[(data.sqft <= 2500) & (data.price <= 8000) & (data.price != 0) & (data.bedrooms <= 4) & (data.bathrooms <= 3)
           & (data.sqft != 0)]

# add squared square footage to the table
squared = data.adj_sqft ** 2
squared.name = 'sqft_squared'

squared_beds = data.bedrooms ** 2
squared_beds.name = 'beds_squared'

data = pd.concat([data, squared, squared_beds], axis=1)
#data = pd.concat([data, squared_beds], axis=1)

# create X and y
feature_cols = ['adj_sqft', 'bedrooms', 'bathrooms', 'sqft_squared', 'beds_squared']

X = data[feature_cols]
y = data.price

# instantiate, fit
lm = LinearRegression()
lm.fit(X, y)

# print coefficients
print("Intercept: %.2f" % lm.intercept_)
# The mean square error
print("Residual sum of squares: %.2f"
      % np.mean((lm.predict(X) - y) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % lm.score(X, y))
print(zip(feature_cols, lm.coef_))

# calculate predictions for the data set and plot errors
predictions = lm.predict(X)
errors = predictions - y
errors.name = 'Error'

# visualize the relationship between the features and the response using scatterplots
errors.sort()
errors.plot(kind='bar').get_xaxis().set_ticks([])


```

```

Intercept: 247.37
Residual sum of squares: 474109.65
Variance score: 0.73
[('adj_sqft', 3.6996192523678673), ('bedrooms', -157.61331809822042), ('bathrooms', 316.48288907402383), ('sqft_squared',
-0.00059060553655249949), ('beds_squared', 53.514005384693448)]

```

Out[249]: []




```
In [250]: import statsmodels.formula.api as sm
result = sm.ols(formula="price ~ adj_sqft + bedrooms + bathrooms + elevation", data=data).fit()
print result.params
print result.summary()
```

```
Intercept    1355.629789
adj_sqft      1.756130
bedrooms      223.018916
bathrooms     277.936411
elevation     -3.520630
dtype: float64
```

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.721
Model:                  OLS      Adj. R-squared:            0.717
Method:                 Least Squares    F-statistic:          192.5
Date:                  Sun, 16 Aug 2015    Prob (F-statistic):    2.69e-81
Time:                  13:17:50    Log-Likelihood:        -2416.8
No. Observations:      303    AIC:                    4844.
Df Residuals:          298    BIC:                    4862.
Df Model:              4
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[95.0% Conf. Int.]	
Intercept	1355.6298	131.263	10.328	0.000	1097.309	1613.950
adj_sqft	1.7561	0.109	16.059	0.000	1.541	1.971
bedrooms	223.0189	75.983	2.935	0.004	73.488	372.550
bathrooms	277.9364	99.297	2.799	0.005	82.524	473.349
elevation	-3.5206	1.182	-2.979	0.003	-5.846	-1.195

```
=====
Omnibus:                 18.909    Durbin-Watson:          1.769
Prob(Omnibus):           0.000    Jarque-Bera (JB):        22.344
Skew:                    0.532    Prob(JB):                1.41e-05
Kurtosis:                 3.800    Cond. No.:               5.15e+03
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 5.15e+03. This might indicate that there are
strong multicollinearity or other numerical problems.
```

```
In [251]: from mpl_toolkits.basemap import Basemap
import fiona
```

In [252]: plt.figure(figsize=(12,12))

```
# Create the Basemap
event_map = Basemap(projection='merc',
                    resolution='h', epsg=2227,
                    lat_0 = 37.7, lon_0=-122.4, # Map center
                    llcrnrlon=-122.55, llcrnrlat=37.7, # Lower left corner
                    urcrnrlon=-122.35, urcrnrlat=37.85) # Upper right corner

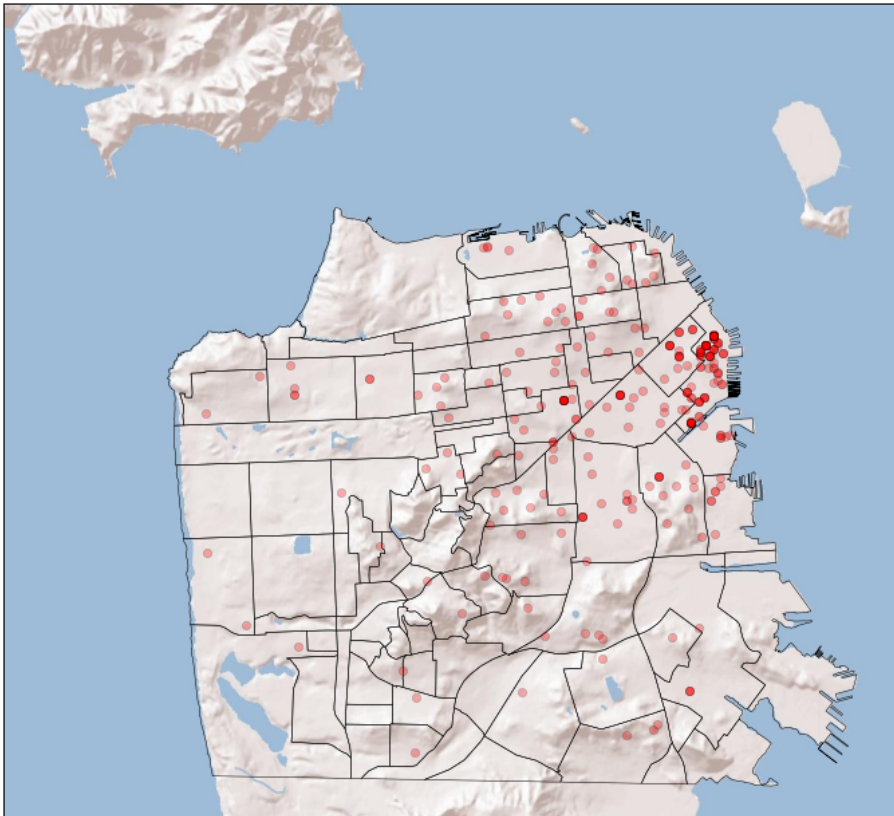
# Draw important features
event_map.arcgisimage(service='World_Shaded_Relief', xpixels = 1500, verbose= True)

# add neighborhoods
event_map.readshapefile(
    'data/Realtor_Neighborhoods_4326/hoods_4326', 'SF', color='black', zorder=2)

# create array storing lats and longs
listing_coords = zip(data.latitude,data.longitude)

# Draw the points on the map:
for longitude, latitude in listing_coords:
    x, y = event_map(latitude, longitude) # Convert lat, long to y,x
    event_map.plot(x,y, 'ro', alpha=0.3)
```



http://server.arcgisonline.com/ArcGIS/rest/services/World_Shaded_Relief/MapServer/export?bbox=5968621.97922,2083843.65958,6027551.68158,2137245.61137&bboxSR=2227&imageSR=2227&size=1500,1359&dpi=96&format=png32&f=image



```
In [253]: price_per_adj_foot = data['price'] / data['adj_sqft']
price_per_adj_foot.name = 'price_per_adj_foot'
adj_price = data['price'] * data['area_multiplier']
adj_price.name = 'adj_price'
data = pd.concat([data, price_per_adj_foot, adj_price], axis=1)

# visualize the relationship between the features and the response using scatterplots
data.plot(kind='scatter', x='elevation', y='adj_price')
data.plot(kind='scatter', x='elevation', y='price_per_adj_foot')
```

Out[253]: <matplotlib.axes._subplots.AxesSubplot at 0x10d822b90>



In []: